



ROI TRAINING
MAXIMIZE YOUR TRAINING INVESTMENT™

Activity: Creating a ClusterIP and Load Balancer

Creating the api-service.yaml

- In Google Cloud Shell, change into your `kubernetes-config` folder:
`cd ~/eventsapp/kubernetes-config`
- In the `kubernetes-config` folder, create a new file named `api-service.yaml`
 - Copy/paste the contents shown here:
- Which pods will be added to this service?

```
apiVersion: v1
kind: Service
metadata:
  labels:
    app: events-api-svc
    name: events-api-svc
spec:
  ports:
    - port: 8082
      protocol: TCP
      targetPort: 8082
  selector:
    app: events-api
    ver: v1.0
  type: ClusterIP
```

Deploying the api-service.yaml

- Deploy the events-api service:
`kubectl apply -f api-service.yaml`
- Verify it was deployed:
`kubectl get service`

```
$ kubectl get service
NAME                TYPE        CLUSTER-IP    EXTERNAL-IP    PORT(S)        AGE
events-api-svc      ClusterIP   10.8.10.141   <none>         8082/TCP       4m1s
kubernetes          ClusterIP   10.8.0.1      <none>         443/TCP       4h49m
```

- Notice the service name is events-api-svc
 - That name will be resolved to the cluster-ip by the kube DNS

Testing a Service DNS Resolution

- You will exec into the events-website pod and connect to the name of the events-api-svc service
 - Retrieve a list of running pods:
`kubectl get pods`
 - Copy the name of the events-web pod
 - Execute into the events-web pod
`kubectl exec -i -t PASTE-POD-NAME -- /bin/bash`
 - Run the following command in the pod:
`curl http://events-api-svc:8082/events`
 - This should connect to the events-api-svc service and return a JSON list of events

Modify the web-deployment.yaml

- Remember, the events-website service uses an environment variable (called SERVER) to know where to connect for the events-api service
 - Environment variables can be set on a Kubernetes deployment
 - In this case, the hostname can be set to the name of the events-api-svc
- Edit your `web-deployment.yaml` file
 - At the end of the file, add the three highlighted lines shown here:
 - Be sure to note the indentation
- Then reapply the file:
`kubectl apply -f web-deployment.yaml`

Add these three lines

```
containers:
- image: IMAGE-URL
  name: events-web
  ports:
  - containerPort: 8080
  env:
  - name: SERVER
    value: "http://events-api-svc:8082"
```

Creating the web-service.yaml

- In Cloud Shell, change into your `kubernetes-config` folder:
`cd ~/eventsapp/kubernetes-config`
- In the `kubernetes-config` folder, create a new file named `web-service.yaml`
 - Copy/paste the contents shown here:

```
apiVersion: v1
kind: Service
metadata:
  labels:
    app: events-web-svc
  name: events-web-svc
spec:
  ports:
    - port: 80
      protocol: TCP
      targetPort: 8080
  selector:
    app: events-web
    ver: v1.0
  type: LoadBalancer
```

Deploying the web-service.yaml

- Deploy the events-web service:

```
kubectl apply -f web-service.yaml
```

- Verify it was deployed:

```
kubectl get service
```

```
$ kubectl get service
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
events-api-svc	ClusterIP	10.8.10.141	<none>	8082/TCP	3h34m
events-web-svc	LoadBalancer	10.8.13.139	34.70.114.185	80:30062/TCP	3h6m
kubernetes	ClusterIP	10.8.0.1	<none>	443/TCP	2m

- If you do not have an EXTERNAL-IP for the events-web-svc, wait a few seconds and try again

Accessing the Load Balancer

- Copy the EXTERNAL-IP for the events-web-svc and visit it in a new browser tab
- The application should now be working again:
 - The app is now running in Kubernetes!
 - Try adding a few events
- If you see a connection error as shown here:
 - The SERVER environment variable was not added to the web-deployment

Doug's Events App

Add an event

Company Pet Show

Super-fun with furry friends!

When: November 6 at Noon

Where: Reston Dog Park



Likes: 0

2021-09-30:12:00

Doug's Events App

Add an event

Error: connect ECONNREFUSED 127.0.0.1:8082

Experiment with Replicas

- Scale the events-web service:
 - Modify the `web-deployment.yaml` to have 3 replicas
 - Apply the file and test the application
 - Everything should still work fine
- Scale the events-api service:
 - Modify the `api-deployment.yaml` to have 3 replicas
 - Apply the file and test the application
 - Now there is a problem with how the events data is stored
 - This is because the events-api service is storing state
 - Each copy has its own state
 - This will be corrected when a database is added

Experiment with Replicas (continued)

- Set the replicas back to 1 in both `web-deployment.yaml` and `api-deployment.yaml`
 - Apply both files
 - Verify the pods have scaled back to 1 each

Clean Up

- There is no need to clean up the deployments, services, or pods
- Leave them running on your cluster

Success

- **Congratulations!** You have successfully created Kubernetes services
 - Created a ClusterIP for the events-api service
 - Created a load balancer for the events-web service
 - Set environment variables for a Kubernetes deployment