



**ROI**TRAINING  
MAXIMIZE YOUR TRAINING INVESTMENT™

# Activity: Network Policies

# Introduction

- In this activity, you will use network policies to
  - Restrict network traffic between pods
  - Isolate traffic based on namespaces
- The activity assumes the Events App case study is deployed to the cluster
  - If it is not, you need to do that first

# Enable Network Policies on the Cluster

- Before performing this lab, your cluster must have network policies enabled
- To create a new GKE cluster with network policies enabled:  

```
gcloud container clusters create "cluster-1" --zone "us-central1-c"  
--machine-type "e2-small" --preemptible --num-nodes "3"  
--enable-network-policy
```
- To update an existing GKE cluster to enable network policies:  

```
gcloud container clusters update cluster-1 --zone us-central1-c  
--enable-network-policy
```

# Verify the App is Still Accessible

- Retrieve the EXTERNAL IP of the web-svc and verify you can reach it from a web browser  
`kubectl get service`
- Keep the browser tab open for future testing

# Create a Namespace

- Create a new namespace called test:  
`kubectl create namespace test`
- List all namespaces:  
`kubectl get namespace`
  - You should see 4 namespaces
- View the pods in the default and test namespaces:  
`kubectl get pods`  
`kubectl get pods -n test`
  - The Events App should be in the default namespace
  - Nothing is in the test namespace

# Testing Network Connectivity

- Run the following command to create a pod in the test namespace and open an interactive shell in the pod

```
kubect1 run net-pol-demo --rm -i -t --image=curlimages/curl -n test -- sh
```

- The --rm option will cause the pod to automatically delete when you exit the shell
- The image used is an official docker image with curl
- Once deployed you should see:

```
If you don't see a command prompt, try pressing enter.  
/ $
```

# Testing Network Connectivity

- At the pod prompt, verify you can get out of the cluster:  
`curl www.google.com`
  - You should see the google web page content returned
- At the pod prompt, verify you can get to the api service on port 8082 in the default namespace:  
`curl events-api-svc.default:8082/events`
  - You should see the events objects returned in JSON
  - The `.default` is because the `events-api-svc` is in the default namespace

# Network Policy to Deny Other Namespaces

- Open a new Cloud Shell tab by clicking the +
- In the new Cloud Shell tab, change into your kubernetes-config folder:  
`cd ~/eventsapp/kubernetes-config`
- In the kubernetes-config folder, create a new file named `deny-from-other-ns.yaml`
  - You can create the file with the Cloud Shell editor
  - Copy the contents from this slide:

```
kind: NetworkPolicy
apiVersion: networking.k8s.io/v1
metadata:
  namespace: default
  name: deny-from-other-namespaces
spec:
  podSelector:
    matchLabels: {}
  ingress:
    - from:
      - podSelector: {}
```



# Network Policy to Deny Other Namespaces

- Apply the network policy:  
`kubectl apply -f deny-from-other-ns.yaml`
- Switch back to the Cloud Shell tab connected to the net-pol-demo pod
  - Try to reach the api service in the default namespace again:  
`curl events-api-svc.default:8082/events`
  - This time the command should just hang and not connect
  - You can no longer connect to anything in the default namespace from outside that namespace
  - Traffic within the default namespace is still allowed
- Switch back to the browser displaying the Events App
  - Try creating a new event
  - This also does not work because you are outside the namespace

# Allow Traffic to the Website Pods

- In the kubernetes-config folder, create a new file named `allow-web-traffic.yaml`
  - Copy the contents from this slide:
- Apply the file when ready
- Test the connectivity again
  - You should now be able to load the events app in the browser
  - Accessing the api-svc from the pod in the test namespace is still blocked

```
kind: NetworkPolicy
apiVersion: networking.k8s.io/v1
metadata:
  namespace: default
  name: allow-web-traffic
spec:
  podSelector:
    matchLabels:
      app: events-web
  ingress:
    - from:
      - ipBlock:
          cidr: '0.0.0.0/0'
```

# Block Egress Traffic

- In the kubernetes-config folder, create a new file named `block-egress.yaml`
  - Copy the contents from this slide:
- Apply the file when ready
- Switch to the Cloud Shell connected to the net-pol-demo pod, and try:  
`curl www.google.com`
  - This will no longer work
  - Even DNS is blocked:

`nslookup www.google.com`

```
kind: NetworkPolicy
apiVersion: networking.k8s.io/v1
metadata:
  name: default-deny-all-egress
  namespace: test
spec:
  policyTypes:
    - Egress
  podSelector: {}
  egress: []
```

# Allow Egress DNS

- In the kubernetes-config folder, create a new file named

`allow-egress-dns.yaml`

- Copy the contents from this slide:
- Apply the file when ready
- Switch to the Cloud Shell connected to the net-pol-demo pod, and try:  
`nslookup www.google.com`

```
kind: NetworkPolicy
apiVersion: networking.k8s.io/v1
metadata:
  name: allow-egress-dns
  namespace: test
spec:
  policyTypes:
  - Egress
  podSelector: {}
  egress:
  - ports:
    - port: 53
      protocol: UDP
    - port: 53
      protocol: TCP
  - to:
    - namespaceSelector: {}
```

# Clean Up

- Switch to the Cloud Shell connected to the net-pol-demo pod
  - Type exit to exit the shell
  - This will automatically delete the pod
- List all network policies:  
`kubectl get netpol`  
`kubectl get netpol -n test`
- Feel free to delete the network policies if you want  
`kubectl delete netpol <Name> [-n test]`

# Success

- **Congratulations!** You have used network policies to implement pod level firewall rules
  - Restricted network traffic between pods
  - Isolated traffic based on namespaces