ROI TRAINING
MAXIMIZE YOUR TRAINING INVESTMENT

Kubernetes Essentials

ROI TRAINING
MAXIMIZE YOUR TRAINING INVESTMENT

The following course materials are copyright protected materials. They may not be reproduced or distributed and may only be used by students attending the *Kubernetes Essentials* course.

Welcome!

- ROI leads the industry in designing and delivering customized technology and management training solutions
  - Our expertise is in supporting large scale training programs for the global banking community; e.g., cloud and new hire training
- Meet your instructor
  - Name
  - Background

Sample Public Cloud Provider Training Offerings

Course Objectives

- Introduction to running Kubernetes clusters in the cloud
- Quickly deploying autoscaled, load-balanced applications to Kubernetes
- Implement a DevOps automated deployment with a CI/CD pipeline
- Gain hands-on experience with Kubernetes

Qwiklabs

- This course leverages Qwiklabs to provide hands-on tutorials
  - You will be provided with Qwiklabs credits
  - We will perform one lab during the course
  - These credits will be available for approximately 2 weeks to perform additional labs after the course
- You will need to create a free Qwiklabs account:
  - Go to https://www.qwiklabs.com/
  - Click the Join button
  - Follow the prompts to create new account
  - You may receive an email to confirm your account
  - Once logged in, click on your profile (top-right corner), and verify you have credits

Agenda

Microservice Architecture

Kubernetes

DevOps Automation (CI/CD)

Microservice Architecture

- Divide a large program into a number of smaller, independent services
  - Each service is programmed, deployed, and run separately
- Advantages of microservices include:
  - Reduced risk when deploying new versions
    - Software can be changed and re-deployed without fear of compromising the core application
  - Services scale independently to optimize use of infrastructure
  - Easier to innovate and add new features
  - Can use different languages and frameworks for different services

Monolithic vs. Microservice Architecture

- Monolithic applications implement all features in a single code base
  - Single database for all data
- Microservices break large programs into a number of smaller services
  - Each service manages its own data

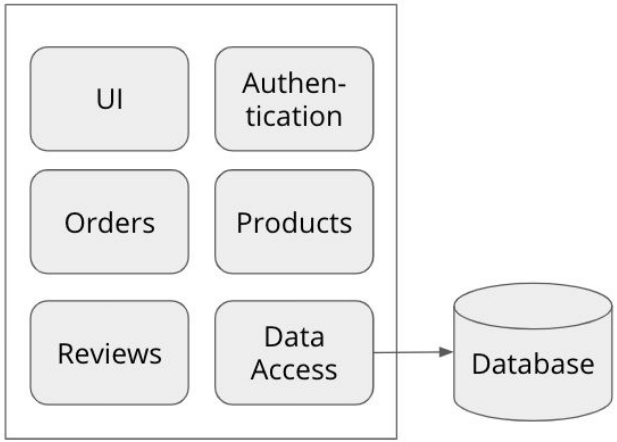

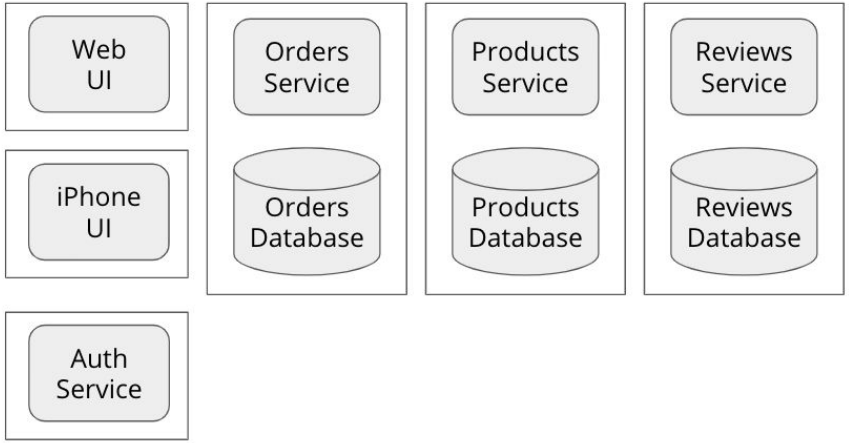

Recognizing Microservice Boundaries

- Defining service boundaries is a matter of decomposition
  - Deciding what pieces of the larger application should be separated
- Typically done by features to minimize dependencies across services
  - Balance service, interest rate service, amortization service, underwriting service, etc.
  - Reviews service, Order Processing service, Product Catalog service, etc.
- Sometimes services are organized by architectural layer
  - Web, iPhone, and Android user interfaces
  - Services that provide access to shared data
- Some services provide shared behavior to the others
  - For example, authentication service

Virtual Machine-Based Virtualization

- Virtual machine-based virtualization is the ability to virtualize hardware
  - Allows resources of a single physical computer to be shared among more than one virtual machine
  - Each virtual machine (VM) has its own virtual set of resources
  - Each VM provides functionality to execute entire operating systems (OS)
- Software that manages this virtualization is known as a hypervisor
  - Type 1 or native hypervisor
  - Type 2 or hosted hypervisor (not discussed here)

Type 1 (Native) Hypervisor

- A type 1 hypervisor runs directly on the machine hardware
  - Each VM runs in complete isolation from other virtual machines
  - Hypervisor exposes a subset of the actual hardware

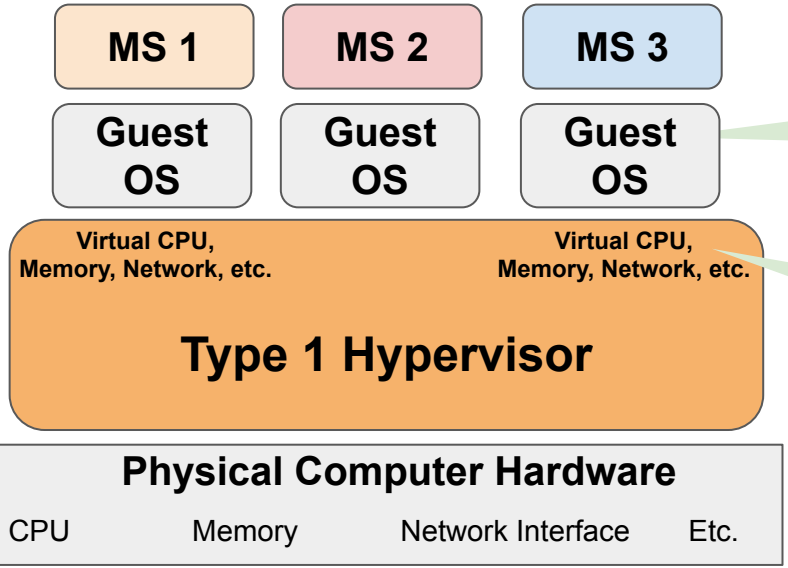

Virtual Machine's Guest OS could be Windows, Linux, etc.

Each Guest OS "sees" a virtualized "view" of hardware

Physical hardware located inside computer

Handling Demand with VMs

- It is standard practice to isolate services on their own virtual machines
  - And run multiple virtual machines for application availability or load
  - Has costs in boot time (minutes) and resources (Gigabytes)

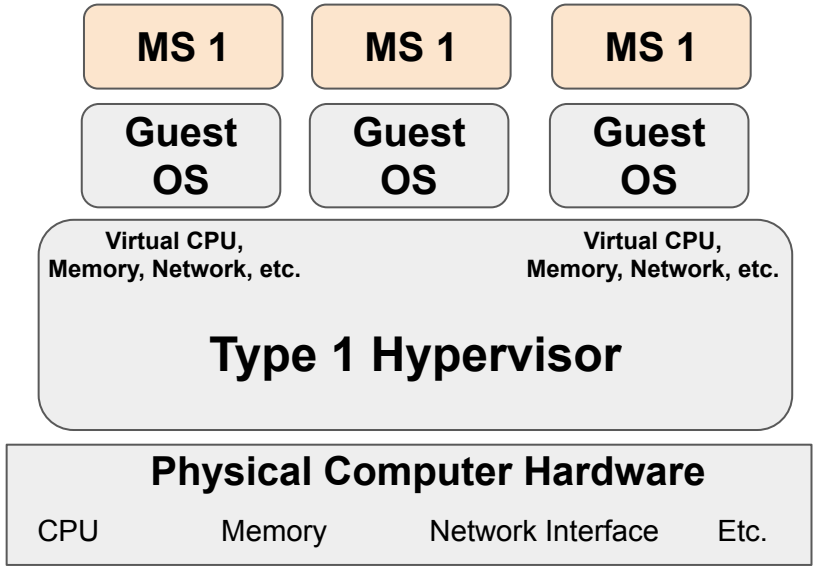

Container-Based Virtualization

- Container-based virtualization is the ability to virtualize the OS kernel
  - Multiple isolated systems, called containers, access a single OS kernel
- Everything required for a piece of software to run is packaged into an isolated container
  - An abstraction layer of the hardware and OS
  - Like an invisible box with configurable access to isolated partitions of file system, RAM, and networking
  - Contain software dependencies, libraries, and settings required to make the software work
- Can deploy and run service without launching an entire VM for each service

Virtual Machines vs. Containers

- Applications hosted on virtual machines require significantly more resources
  - Each application on a VM requires its own operating system
  - In essence, container engines are virtualizing the OS

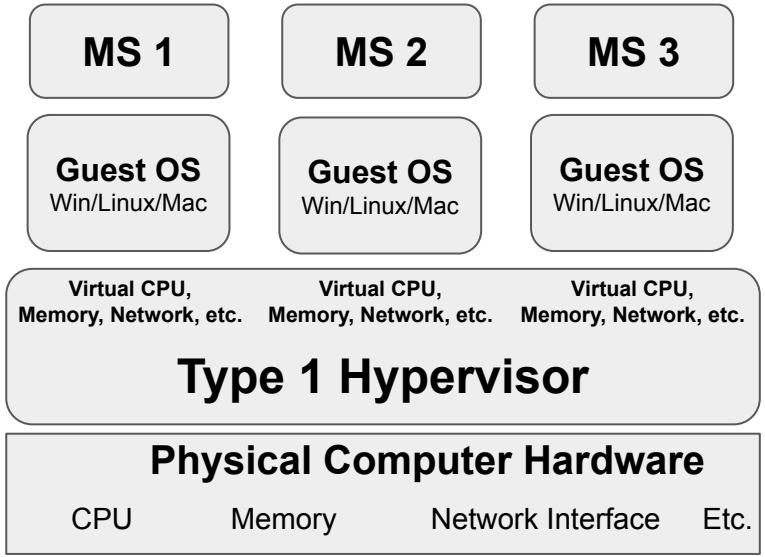

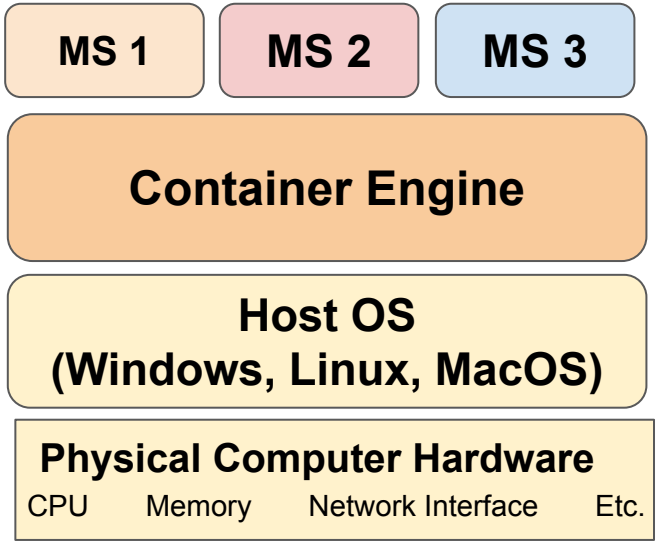

Advantages of Containers

- More efficient
  - Require less resources, less overhead
- Better performance
  - Containers can be very fast to start
    - No VM to boot, which takes tens of seconds or even minutes
- Provides an agile environment
  - Can improve portability across systems
- Encapsulates application dependencies
  - Facilitates a microservices/DevOps approach
  - Works well with a continuous integration strategy

Docker

- Docker is one of the most popular container environments
  - www.docker.com
  - Docker Engine client runs natively on Linux, MacOS, and Windows

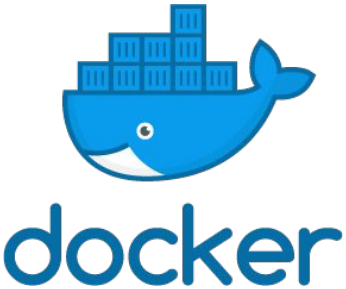

Agenda

Microservice Architecture

Kubernetes

DevOps Automation (CI/CD)

Kubernetes

- Open-source container orchestration system
  - Extremely popular and active open-source project
  - See: https://kubernetes.io/
- Originally developed by Google to run Google's data centers
  - Designed to operate at Google scale
  - Proven and tested running Google's applications
- Wide support
  - Used on Google Cloud when using Google Kubernetes Engine (GKE)
  - Supported by Microsoft in Azure Container Service (AKS)
  - Supported by AWS with Elastic Kubernetes Service (EKS)
  - Red Hat OpenShift
  - And more …

Kubernetes Clusters

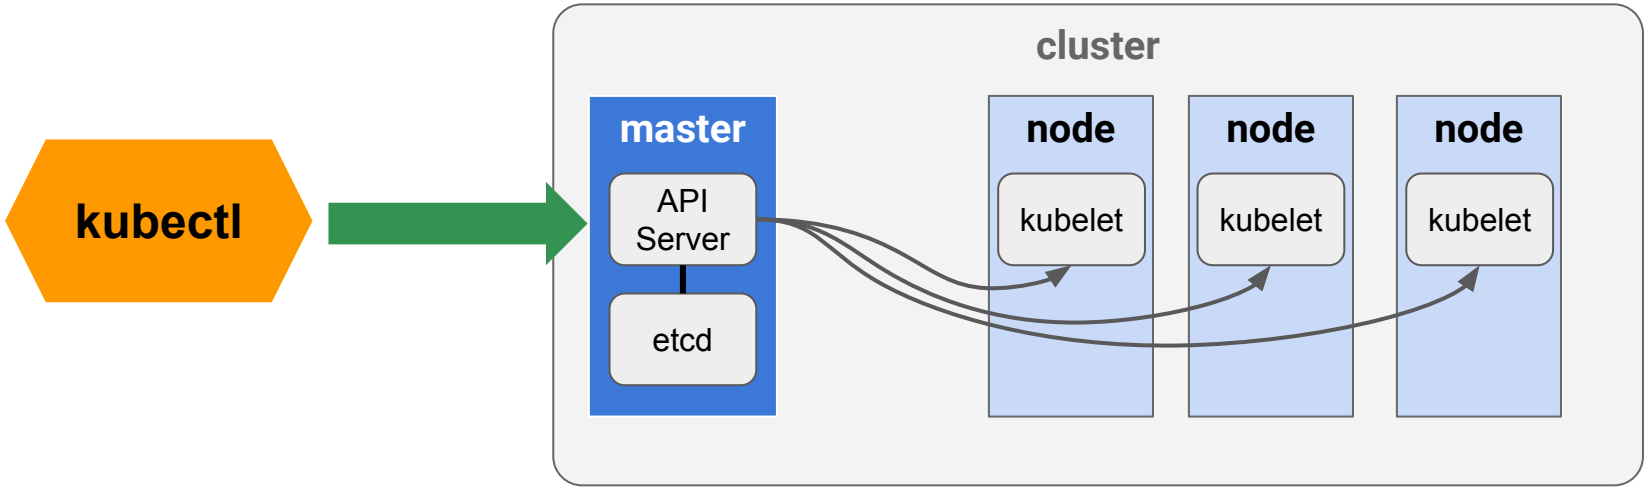

Managed Kubernetes Services

- Many cloud platforms provide a managed Kubernetes service
  - Automates cluster creation and maintenance
  - Clusters are implemented as a collection of cloud VMs
  - Very simple creation using the web console or command line

Tutorial: Getting Started with AWS Elastic Kubernetes Service (EKS)

- https://docs.aws.amazon.com/eks/latest/userguide/getting-started-eksctl.html

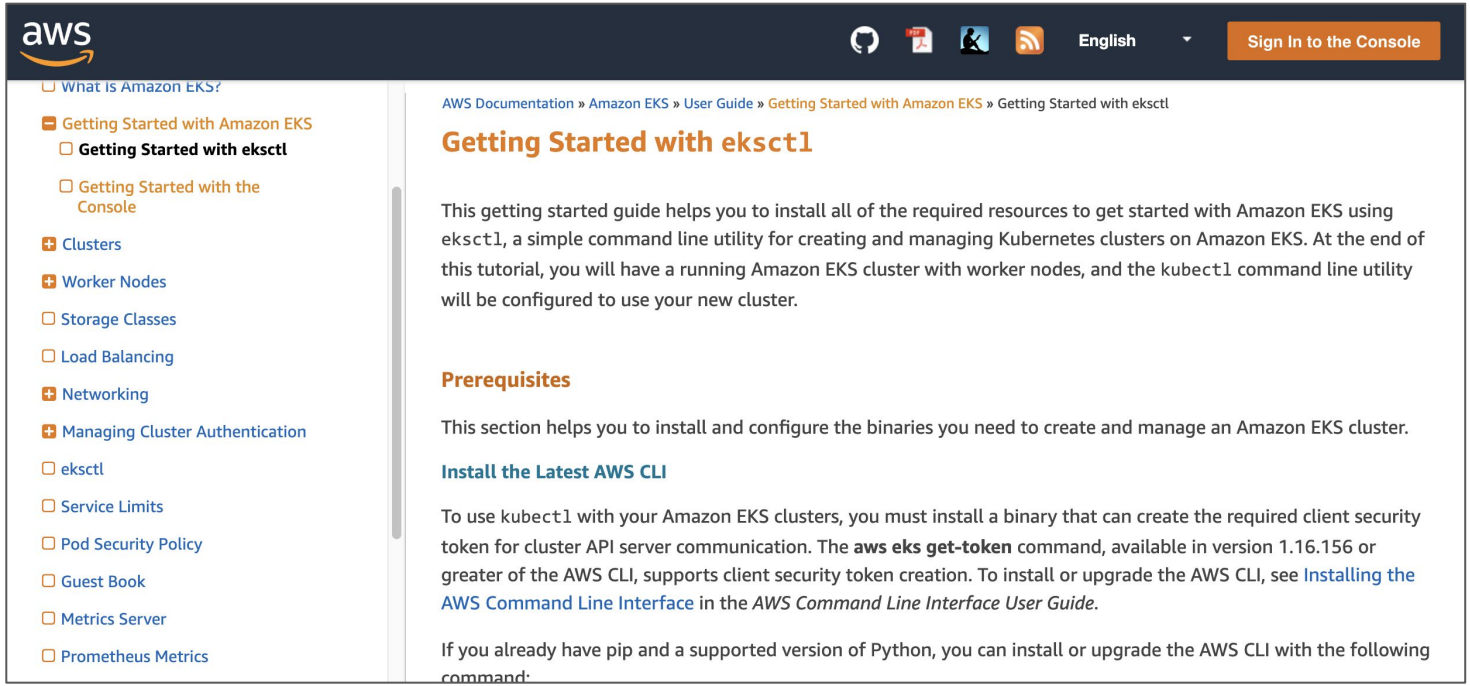

Tutorial: Getting Started with Google Kubernetes Engine (GKE)

- https://cloud.google.com/kubernetes-engine/docs/tutorials/hello-app

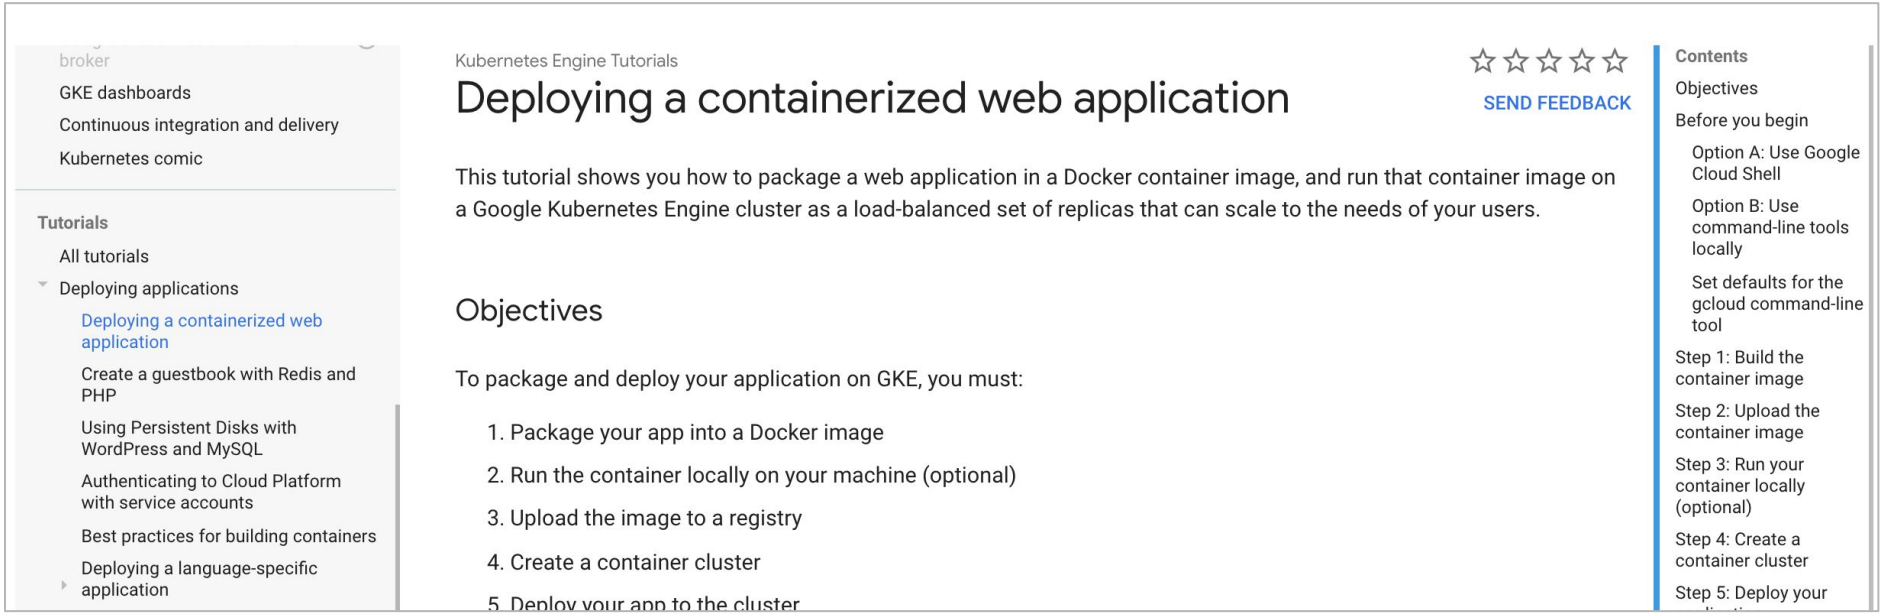

Tutorial: Getting Started with Azure Kubernetes Service (AKS)

- https://docs.microsoft.com/en-us/azure/aks/

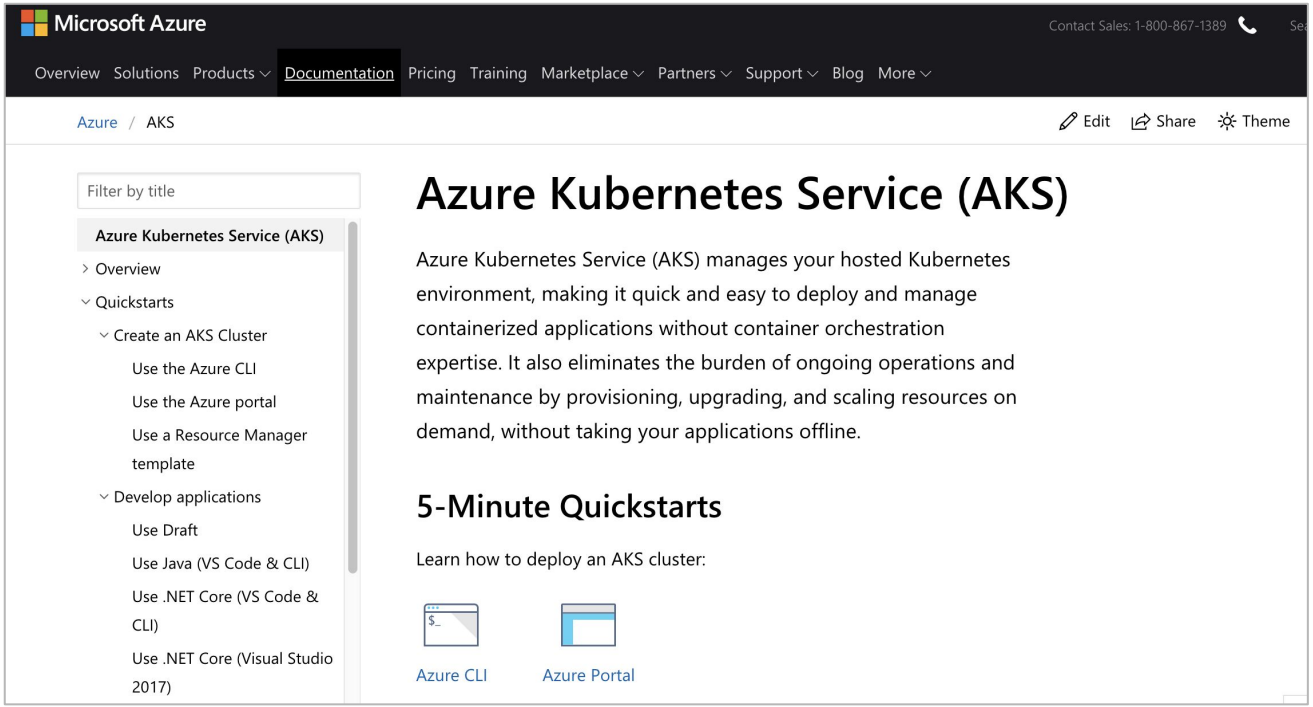

Kubernetes Terms

- Clusters are collections of machines that will run containers
  - Each machine is a node in the cluster
- Pods are the smallest unit of deployment
  - Usually pods represent a single container
- Deployments are configurations that define service resources
- Replication controllers are used to create multiple instances of a pod
  - Guarantee pods are healthy and the right number exist
- Load balancers route requests to pods
- Autoscalers monitor load and turn pods on or off

Kubernetes Components

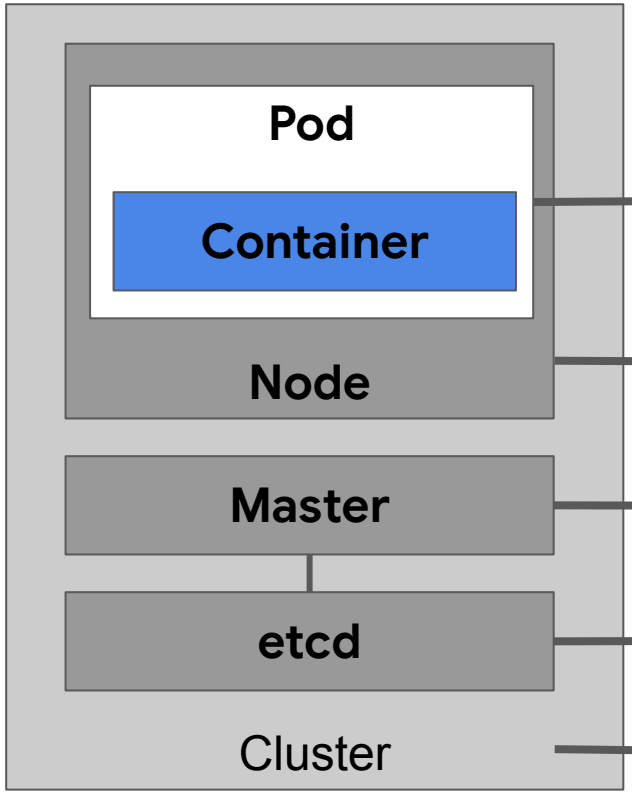

A **pod** is a collection of one or more **containers**, deployed on nodes

Nodes are machines that run the containers, talk to the master

The **master** controls the cluster

etcd is for storing cluster state

A **cluster** is a set of nodes

Kubernetes CLI

- Kubernetes is automated with a combination of CLI commands and configuration code
  - `kubectl` is the Kubernetes CLI

kubectl Command	Description
kubectl get nodes	Show information about the cluster nodes.
kubectl apply -f <file-name>	Apply a configuration file
kubectl get pods	Show the running pods.
kubectl get services	Show running services along with addresses and ports.
kubectl delete [name] kubectl delete deployments/spaceinvaders	Delete specified resources.
kubectl exec [options] [pod-name] kubectl exec -it spaceinvaders-55c88695bb-bcxb2 -- /bin/bash	Execute a command in the container (the example runs a bash shell).

Deployments Combine Pods with Replica Sets

```
apiVersion: apps/v1beta1
kind: Deployment
metadata:
  name: devops-deployment
  labels:
    <Some code omitted to save space>
spec:
  replicas: 3
  selector:
    <Some code omitted to save space>
  template:
    <Some code omitted to save space>
    spec:
      containers:
      - name: devops-demo
        image: drehnstrom/devops-demo:latest
        ports:
        - containerPort: 8080
```

Kind of resource: Deployment

We want 3 replicas of the pod

This spec defines the pod. The same as the Pod configuration.

Kubernetes Deployments

- Use kubectl to send the config file to the master
  - The master then decides how to deploy the pods

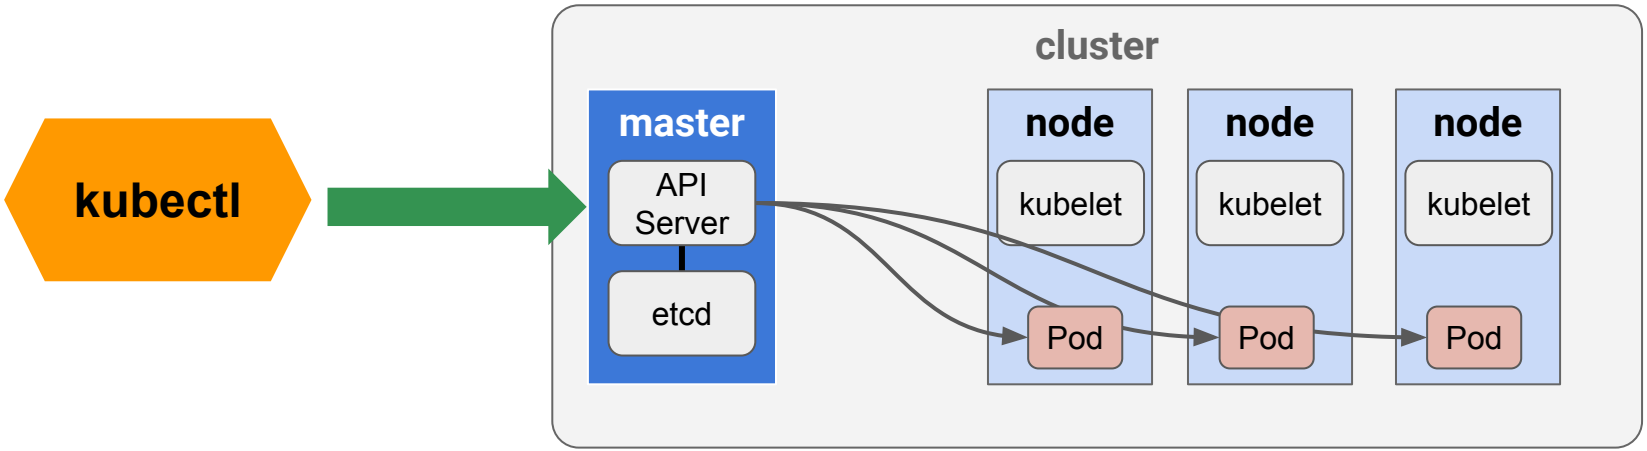

Kubernetes Replica Set

- A replica set ensures the correct number of pods is running
  - And replaces any pod that fails

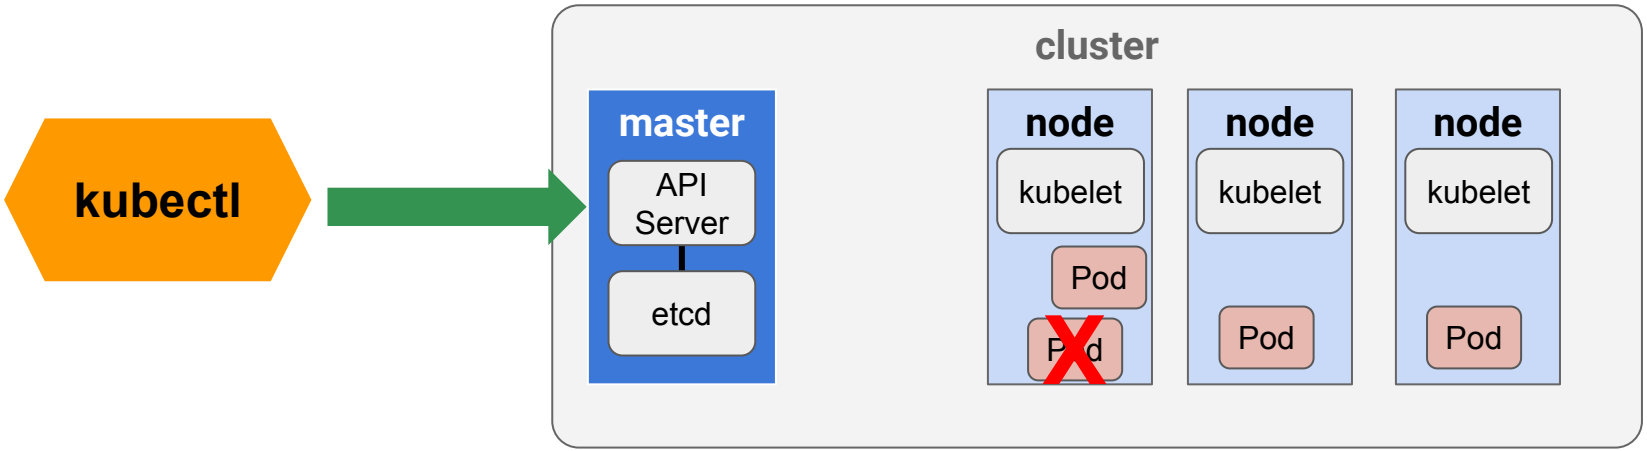

Creating a Deployment from a Configuration File

- Deploy a service based on a configuration file

```
kubectl apply -f kubernetes-config.yaml
```

- Show the running pods

```
kubectl get pods
```

- Show all the deployments

```
kubectl get deployments
```

- Show details of a deployment

```
kubectl describe deployments devops-deployment
```

Adding the Load Balancer to Configuration

```
apiVersion: v1
kind: Service
metadata:
  name: devops-loadbalancer
  labels:
    app: devops
    tier: frontend
spec:
  type: LoadBalancer
  ports:
  - port: 80
    targetPort: 8080
  selector:
    app: devops
    tier: frontend
```

LoadBalancer listens on port 80

Send traffic to pod (container) on port 8080

Route traffic only to pods with these labels

Types of Services

ClusterIP	The default service type. Has only an internal IP address that is only accessible by other services running inside the cluster.
LoadBalancer	A service that provides an external IP address. In Google Cloud, this is implemented as a TCP load balancer. In AWS, this is implemented as an Elastic Load balancer. Not all Kubernetes deployments would support this type of service. Can be expensive if you have lots of services, which means lots of load balancers.
NodePort	Assigns a port between 30000 and 32767 to nodes in your cluster. When a node is accessed at that port, it routes to your service.

Types of Services

- A single application can have multiple services

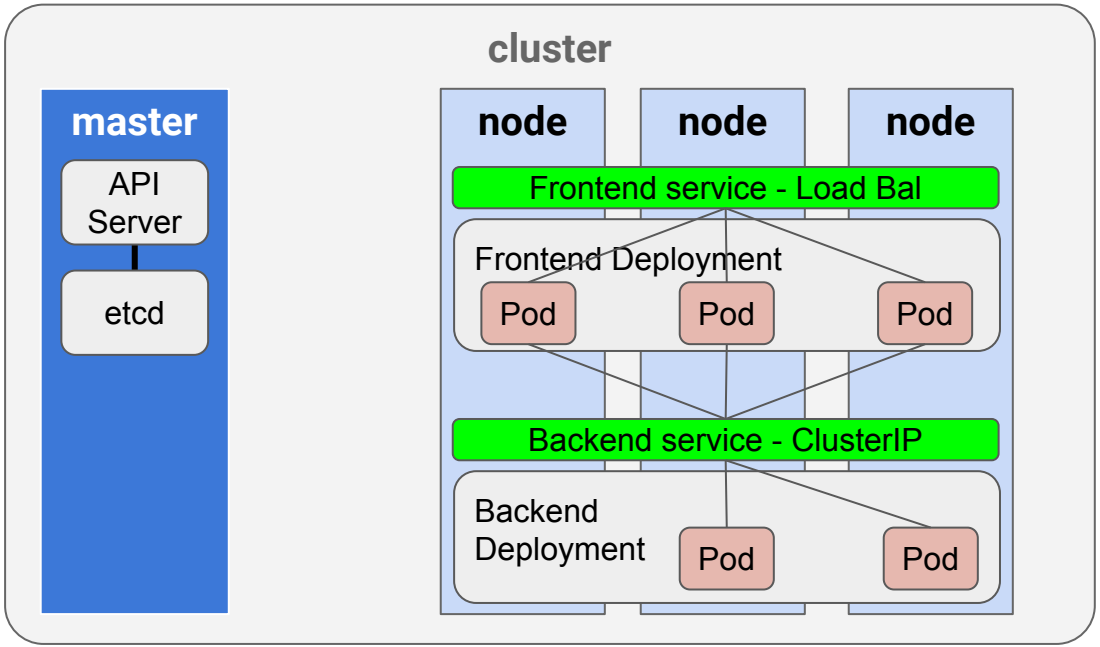

Adding the Autoscaler to Configuration

- To dynamically scale up and down, create an autoscaler
  - Specify min and max number of pods and some metric to monitor

```
apiVersion: autoscaling/v1
kind: HorizontalPodAutoscaler
metadata:
  name: devops-autoscaler
spec:
  scaleTargetRef:
    apiVersion: apps/v1beta1
    kind: Deployment
    name: devops-deployment
  minReplicas: 3
  maxReplicas: 10
  metrics:
  - type: Resource
    resource:
      name: cpu
      targetAverageUtilization: 60
```

Demo: Kubernetes in Action

- Your instructor will demo Kubernetes in action

Kubernetes Engine: Qwik Start

- If you have not already, create a free Qwiklabs account:
  - Go to https://www.qwiklabs.com/
  - Click the Join button
  - Follow the prompts to create new account
  - You may receive an email to confirm your account
  - Once logged in, click on your profile (top-right corner), and verify you have credits
- Perform the Kubernetes Engine: Qwik Start lab:
  - https://www.qwiklabs.com/focuses/878?parent=catalog

Rolling Updates

- Cloud services typically have multiple instances behind a load balancer
  - Multiple replicas of a pod in a Kubernetes cluster
- Rolling deployments update instances incrementally
  - One at a time, 10% at a time, etc.
  - Allows services to be updated with no downtime
- Supported by managed instance groups
- Supported by Kubernetes using the apply command
  - Simply change the container image and re-apply the configuration
  - Can also roll back the update

Rolling Update with Kubernetes

- Initially a deployment creates a single ReplicaSet

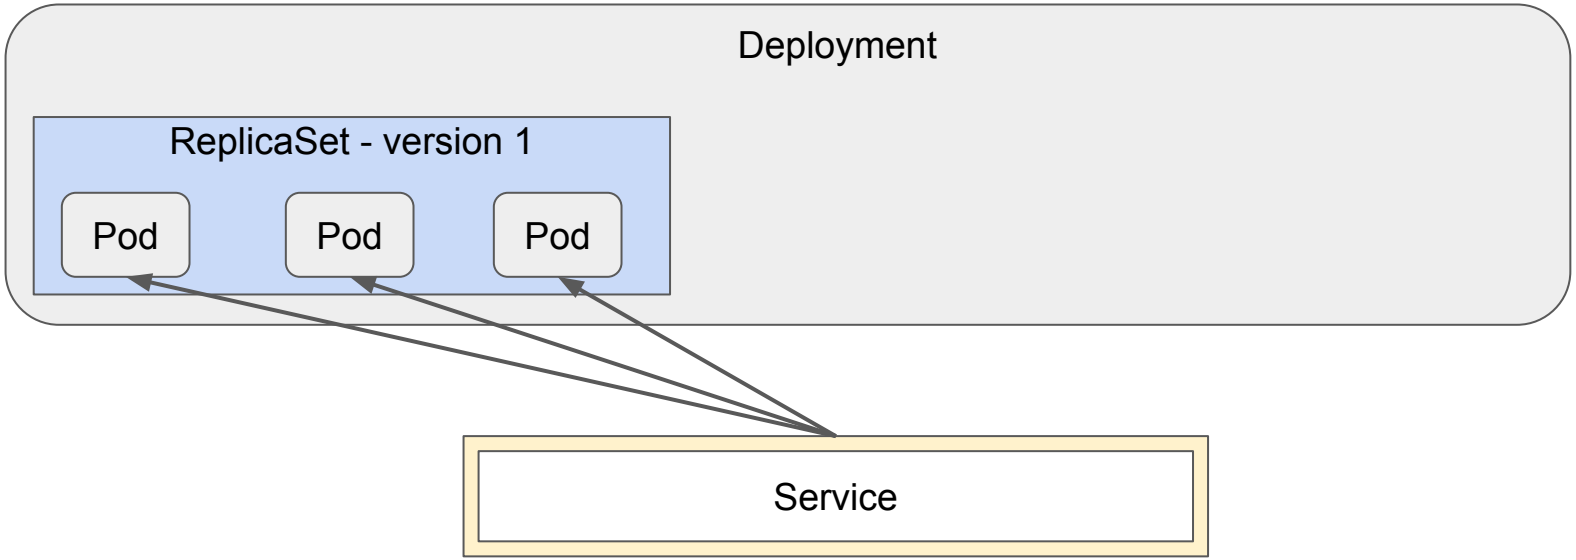

Rolling Update with Kubernetes

- When a container update is applied
  - A new ReplicaSet is created in the same deployment

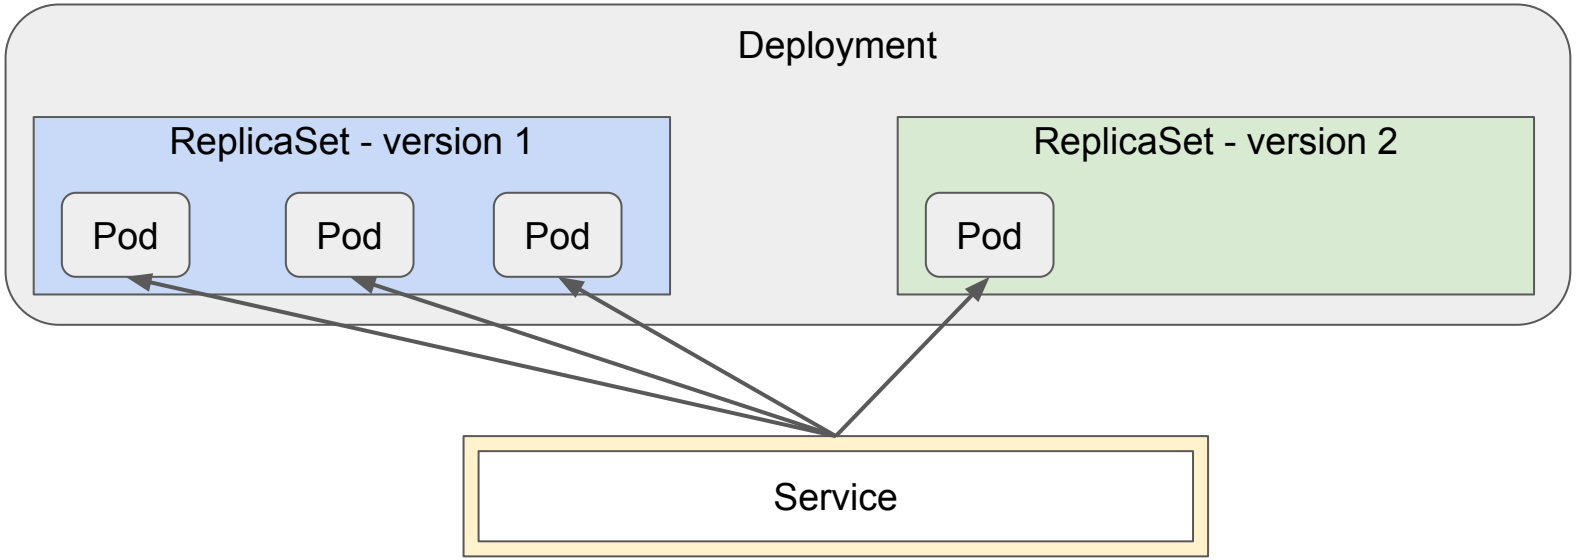

Rolling Update with Kubernetes

- A pod in the old ReplicaSet is deleted

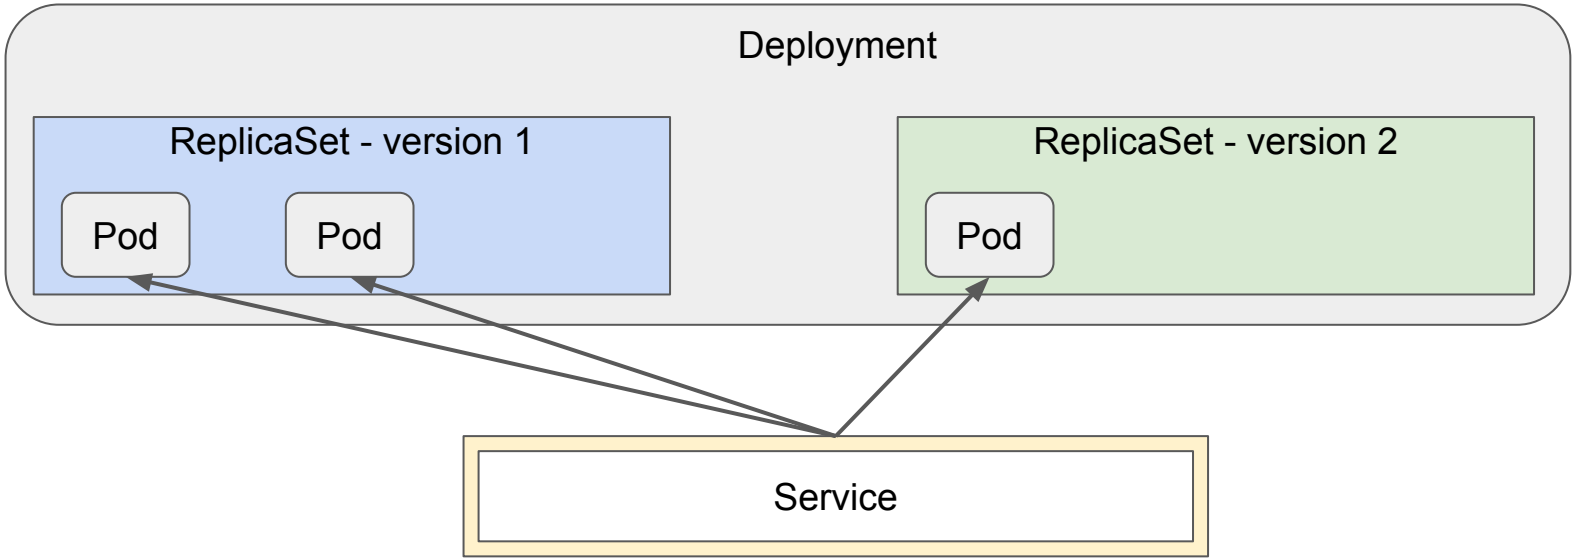

Rolling Update with Kubernetes

- This is repeated

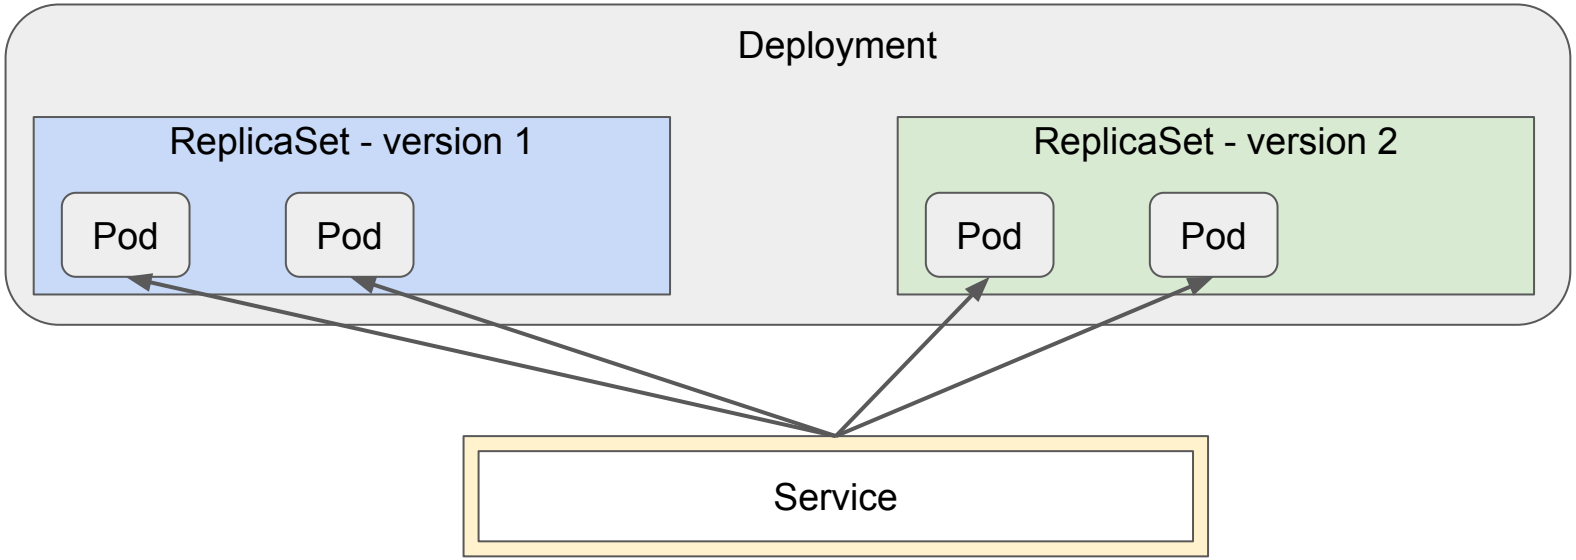

Rolling Update with Kubernetes

- This is repeated

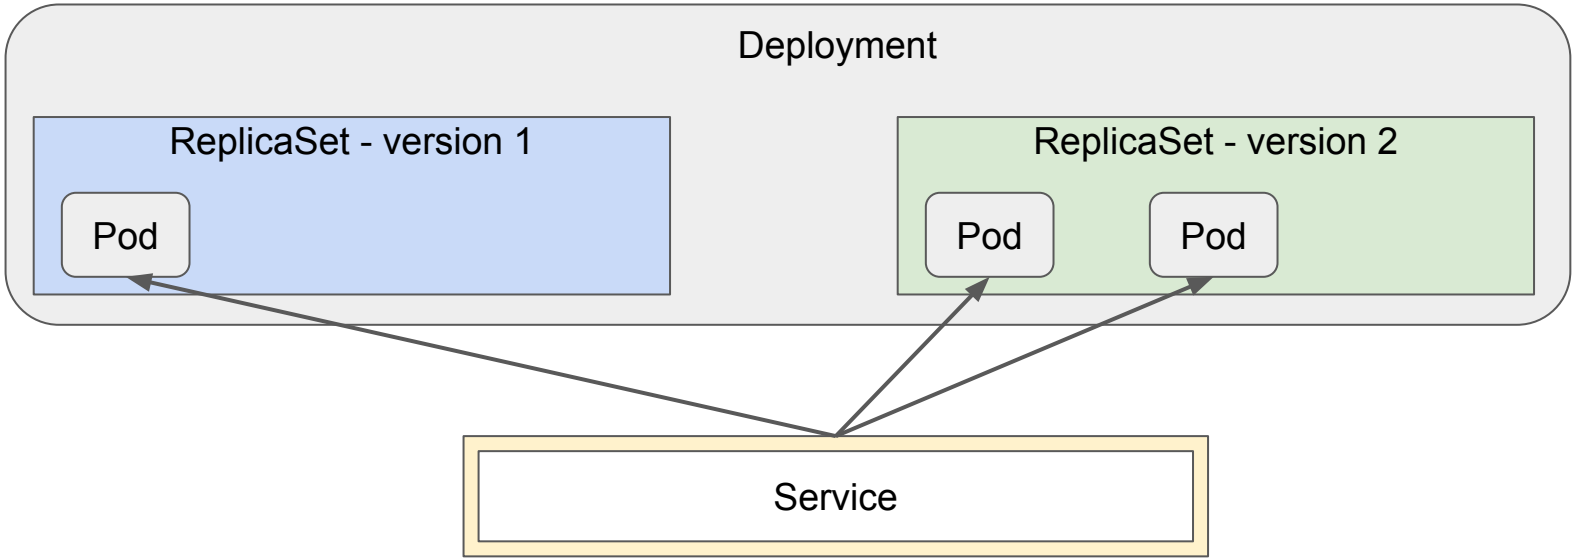

Rolling Update with Kubernetes

- Until all pods have been updated

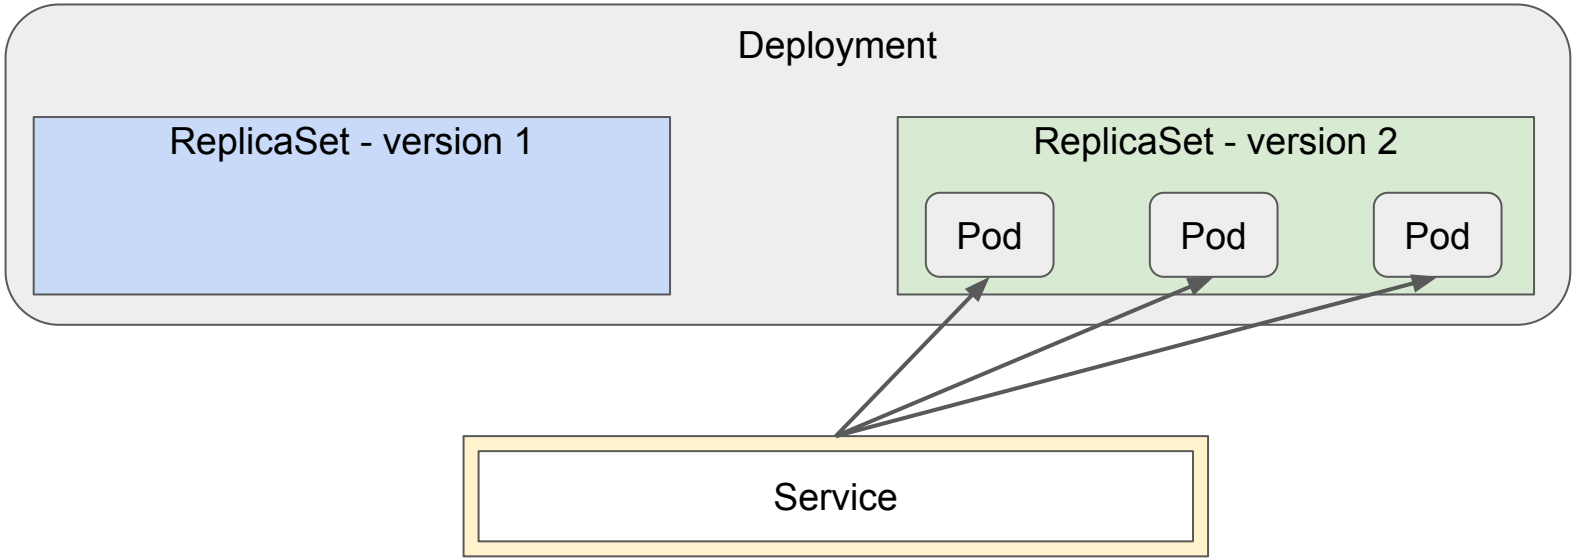

Blue/Green Deployment

- Allows new revisions to be deployed with less risk and no downtime
- There are two copies of the production environment
  - Blue environment is taking requests
  - Green environment is idle
- When deploying a new version:
  - Update new version to green environment leaving the blue environment in place
  - Test the green environment
  - When testing is complete, move the workload to the green environment
  - Green is now blue; can turn the old environment off
- Blue/green deployments also make rolling back to old versions easy

Blue/Green Deployments in Kubernetes

- Service routes all traffic to one version

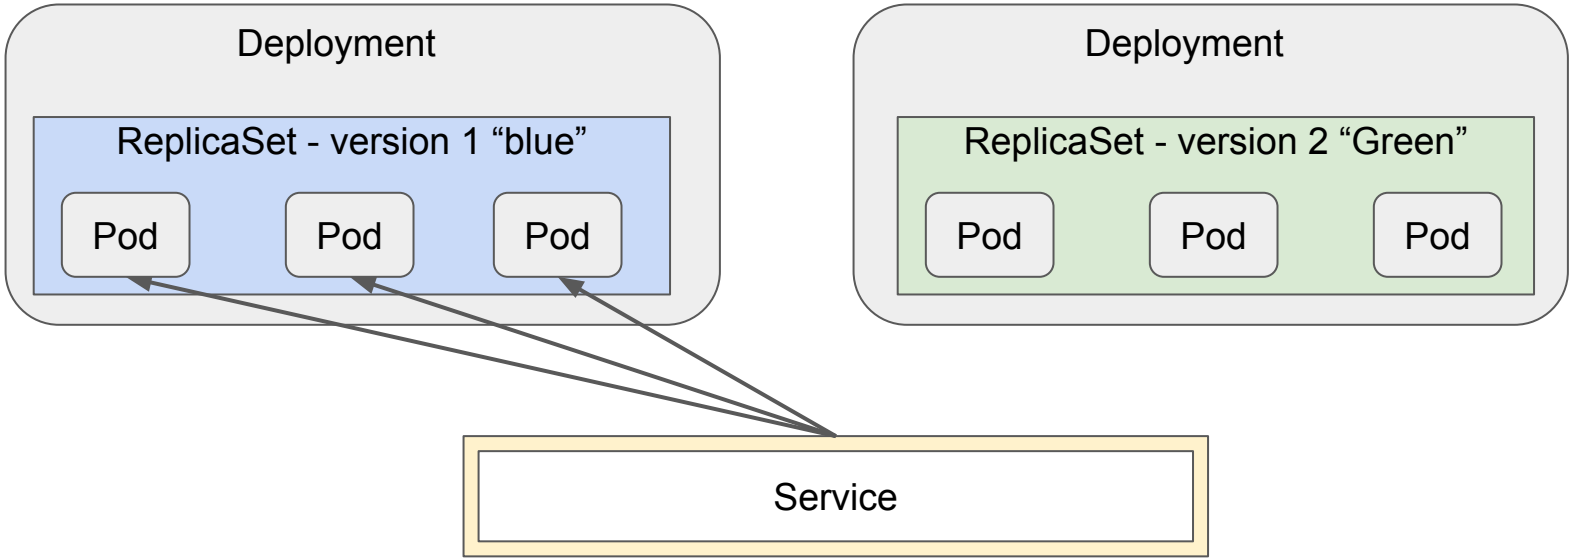

Blue/Green Deployments in Kubernetes

- Can switch the version quickly

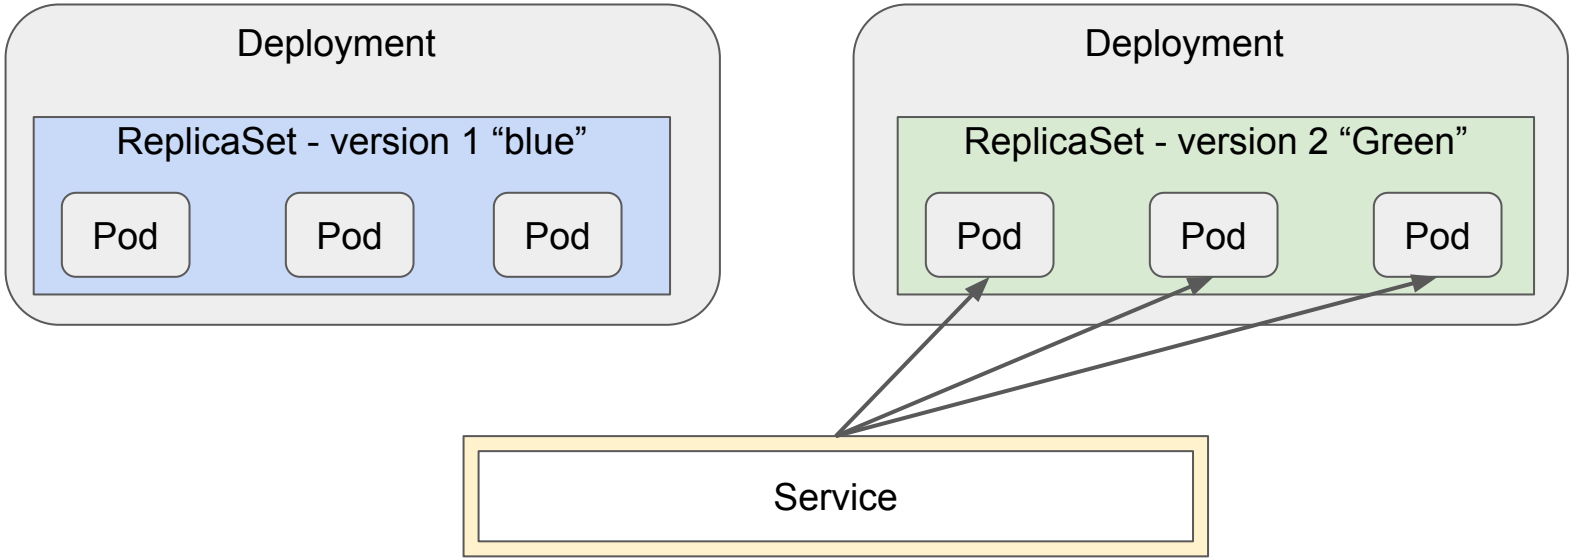

Canary Releases

- A new version of a service is put into production alongside old versions
  - A small subset of select traffic is routed to the canary release
- Canary releases help developers know how a new version will perform
- Canary releases are easy to pull back if they fail their testing

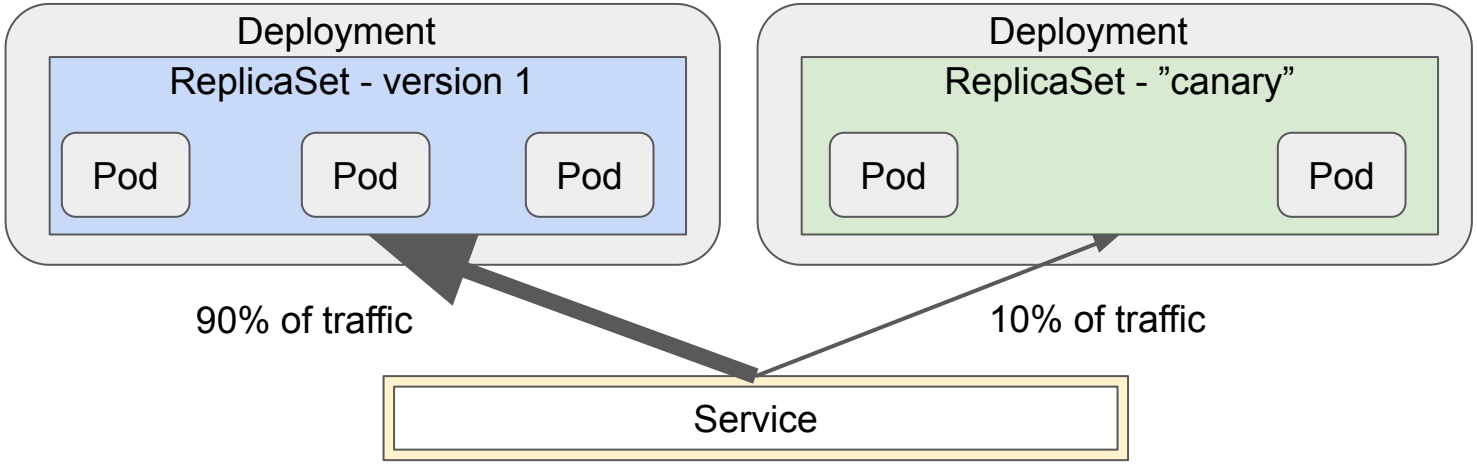

Canary Release in Kubernetes

```
apiVersion: v1
kind: Service
metadata:
  name: devops-loadbalancer
  labels:
    app: devops
    tier: frontend
spec:
  type: LoadBalancer
  ports:
  - port: 80
    targetPort: 8080
  selector:
    app: devops
    tier: frontend
```

1. Services use selectors to determine what pods to route traffic to.
2. Create a new deployment with a new container, but use the same labels as the current deployment.
3. If the current deployment has 3 replicas and the new deployment has 1, then the new deployment gets about 25% of the traffic.

More advanced control possible with an Ingress controller or a service mesh (Istio)

- Outside the scope of this discussion

Demo: Blue/Green Deployments

- Your instructor will demo Blue/Green deployments

Agenda

Microservice Architecture

Kubernetes

DevOps Automation (CI/CD)

Continuous Integration/Continuous Deployment

- CI/CD bridges the gaps between development and operation activities
  - CI = Automated code build and test sequence
  - CD = Automating the entire software release process
  - Enforced automation in building, testing and deployment of services
- Pipelines are used to manage the deployment process
  - Steps are defined within the pipeline
  - Some quality metric is used to trigger the next step
  - The final step is deploying a new release into production
- CI/CD pipelines form the backbone of modern day DevOps

Demo: Automated Build and Deployment

- Your instructor will demo automated builds and deployment

Additional Labs for Homework

- You have been provided with Qwiklab credits
  - Credits will be available for approximately 2 weeks
- Perform one or both of the following labs for more hands on:
  - Introduction to Docker (to gain experience with containers)
    - https://www.qwiklabs.com/focuses/1029?parent=catalog
  - Orchestrating the Cloud with Kubernetes
    - https://www.qwiklabs.com/focuses/557?parent=catalog

Course Summary

In this course, we have:

- Created a Kubernetes cluster in the cloud and deployed a simple workload
- Quickly deployed autoscaled, load-balanced applications to Kubernetes
- Implemented a DevOps automated deployment with a CI/CD pipeline
- Gained hands-on experience with Kubernetes

Cloud Adoption—What's Next?

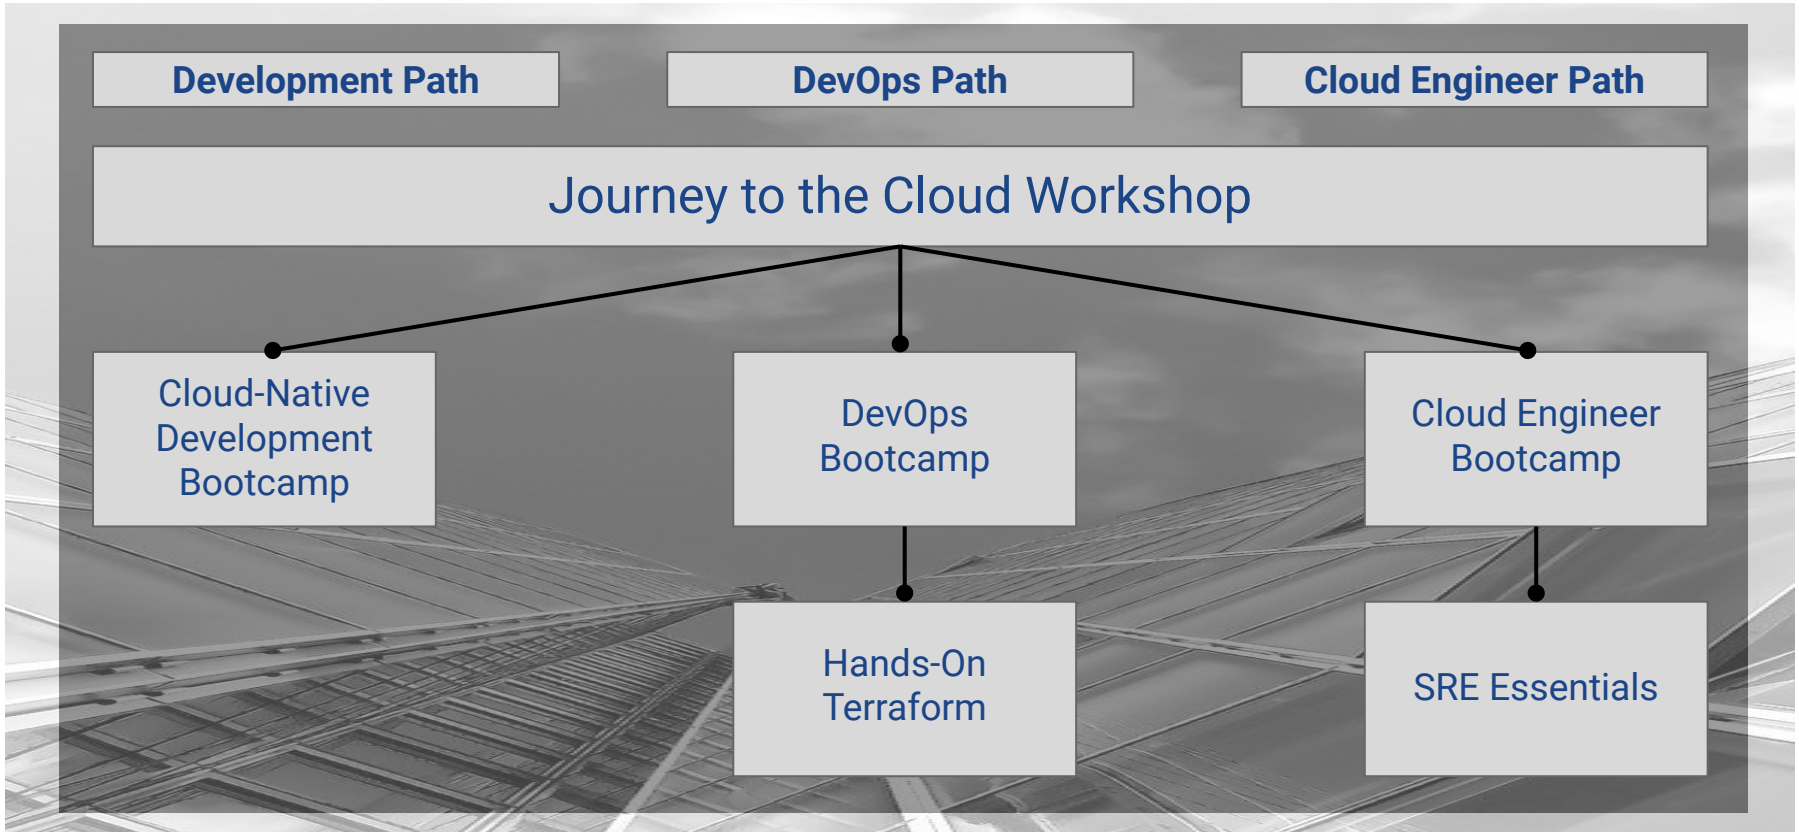

Banking Customer Case Study: Cloud Migration

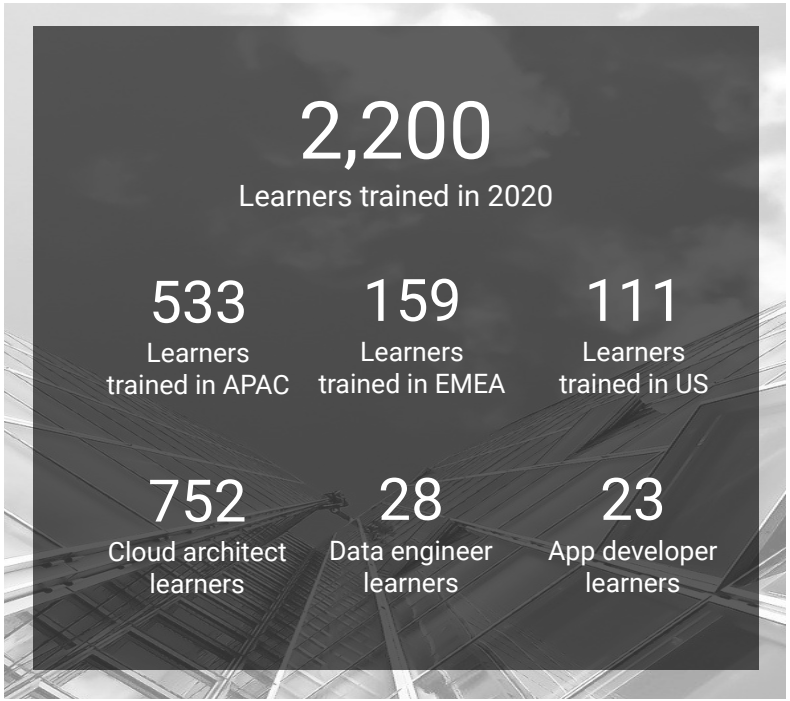

DESIGN PROJECT REVIEWS

83 design review sessions delivered

27 review groups ran concurrently in 1 week and in 4 time zones

5:1 student to instructor ratio

SCALABILITY

2 ROI trainers dedicated to program

4 bootcamps ran concurrently in 1 week

125 trained in 2 time zones in 1 week

BOOTCAMPS BY PATHWAY

70 days of bootcamps

30 Cloud Architect bootcamps

1 Cloud Developer bootcamp

1 Cloud Data Engineer bootcamp

WHAT LEARNERS ARE SAYING

4.66	4.66	4.46	4.73
Overall, this training was worth my time	I would recommend this training to others	I will be able to apply what I've learned to my work	The technical knowledge of the trainer was satisfactory

Our Snapshot

130,000
Trained on cloud technologies and best practices

130+
Certifications held across major public cloud providers

100+
Cloud instructors across NORTHAM, EMEA, and APAC

50+
Google Cloud Authorized trainers (largest team globally)

- Author of AWS, Google Cloud, and Azure courseware
- In-person and virtual delivery options
- 3X Google Cloud Training Partner of the Year (2017, 2018, 2019)
- Global delivery capabilities

ROI Contacts

- **Steve Lockwood**
  steve@lochnetsystems.com
  https://www.linkedin.com/in/steve-lockwood/

- **Mike Prensky**
  Director, Business Development
  mike.prensky@roitraining.com
  781-254-7844